

Lecture 1

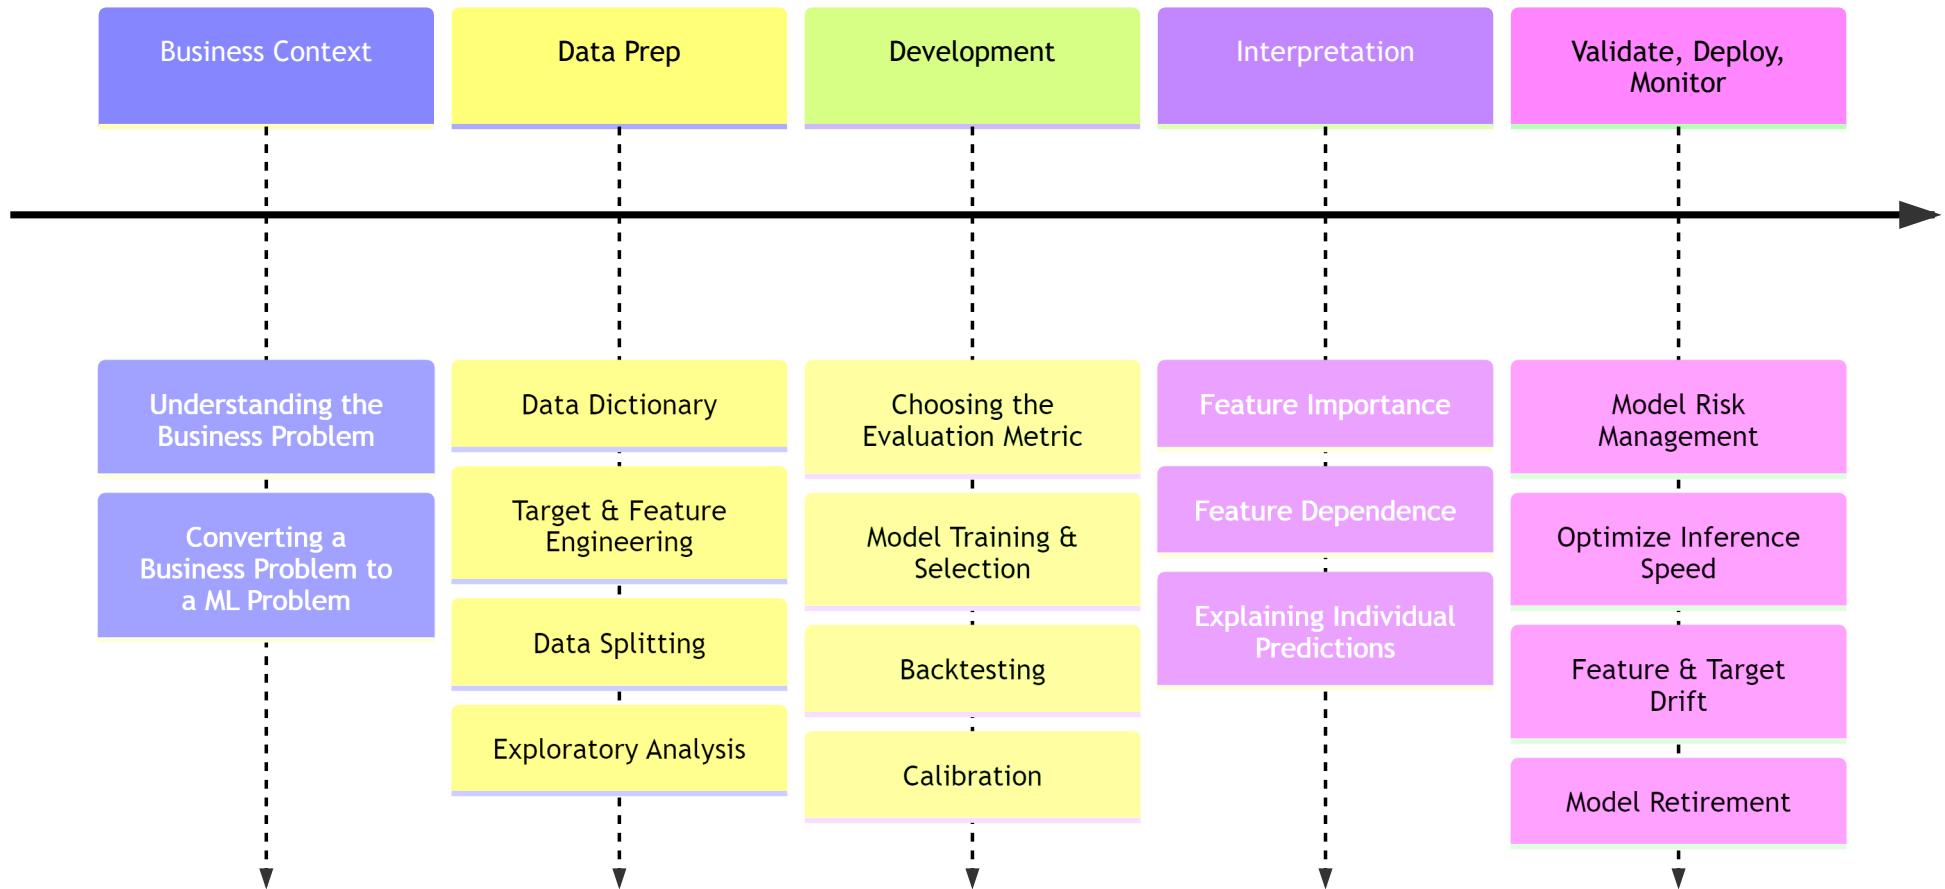
Interpretable Machine Learning for Finance

Instructor

William Chiu, FRM

- Financial Risk Manager (FRM) Certified since 2012
- Northwestern MLDS Alum (Class of 2013)
- Senior Credit Risk Analyst at a credit bureau (2013-2014)
- Head of Statistical Modeling at a bank (2014-2023)
- VP of Predictive Modeling at an insurance company since 2023

Model Lifecycle



Business Context

1. Acquire new clients with a marketing response model (aka prescreen model)
2. Approve (or reject) loan applications with a credit scorecard model
3. Explain rejected loan applications with counterfactuals
4. Estimate lifetime credit losses with a transition matrix model
5. Detect fraud with a cost-sensitive model



Data Sets

1. Borrower data from an Indian Bank and Credit Bureau. The instructor modified the data to include `campaign_id`, `control_group`, and `response_flag`. ([Azad 2024](#)).
2. Loan Performance data from LendingClub. Includes both accepted and rejected loan applications. ([George 2019](#))
3. Single-Family Loan Performance data from Fannie Mae. ([Fannie Mae 2024](#))
4. Simulated data from the *Fraud Detection Handbook* . ([Le Borgne et al. 2022](#))

Model Development Tools



AutoGluon



- AutoGluon is an open-source AutoML framework designed to automate model training and ensembling.
- AutoGluon achieves strong performance by using ensembling techniques and a comprehensive library of preset hyperparameters for each base learner.
- AutoGluon offers data preprocessing like automatic encoding of categorical variables and automatic removal of features with no variance.
- Scikit-learn provides calibration methods (e.g., Platt scaling, isotonic regression) to improve the reliability of predicted probabilities.

Interpretable ML Tools

- Rank features from most to least important
 - Permutation importance
 - SHAP importance
 - Dealing with multicollinear features
- Describe the dependence between features and predictions
 - Partial dependence plots (PDP)
 - Accumulated local effects (ALE) plots
 - SHAP dependence plots



(More) Interpretable ML tools

- Explain individual predictions
 - SHAP values
 - Counterfactuals

Prerequisites

- Comfort using supervised learning methods:
 - Generalized Linear Models
 - Trees with bagging and/or boosting
 - Cross Validation and Bias-Variance Tradeoff
- Familiarity with Python and data science libraries (i.e., pandas and scikit-learn)
- Practice with SQL databases

Assignments and Grading

1. Problem Sets (70%)

- Weekly conceptual questions and programming exercises based on each week's topics.
- Reinforce understanding of weekly material and practical application of concepts.

2. Kaggle Competition (30%)

- Develop a scorecard model for loan underwriting.
- Students will be asked why the model rejected a loan application.
- Grades will depend on model performance in an unlabeled test set and the quality of the explanation.

Introduction to Retail Banking



What is a Retail Bank?

- A retail bank serves individual consumers, households, and small businesses
- Offers deposit and lending products: checking/savings accounts, certificates of deposit, credit cards, and mortgages
- Large retail banks also offer brokerage and wealth management services

Revenue Sources

- Interest collected from loans (e.g., credit cards, auto loans, personal loans, mortgages)
- Transaction fees, commissions, and service charges
- Major contributor to profitability: spread between interest collected from loans and interest paid to depositors

Expense Sources

- Employee compensation
- Interest paid to depositors
- Provision for loan losses

How could a Bank fail?

- Insufficient liquidity (not enough cash to meet obligations)
- Insufficient capital (not enough equity to absorb losses)
- High levels of non-performing loans or significant credit losses
- Sudden loss of depositor confidence leading to a bank run

What is FDIC insurance?

- Protects depositors by guaranteeing their funds up to a certain limit if a bank fails
- Aims to maintain public confidence and prevent bank runs
- Standard insurance amount in the US: \$250,000 per depositor, per insured bank

How do Banks use machine learning?

- Improve the effectiveness of marketing campaigns
- Measure the credit risk of loan applicants
- Estimate the lifetime credit losses of a loan portfolio
- Detect fraudulent activities

Introduction to Credit Bureaus



US Credit Bureaus

- The US has three major credit bureaus:
 - Experian
 - TransUnion
 - Equifax
- Credit bureaus serve three major functions:
 - Collect and maintain consumer credit information
 - Provide credit reports to lenders
 - Generate credit scores

Bureau Data Structure

- **Consumer Information:**
 - Borrower details: Name, address, SSN, birthdate
 - Employment data (if reported)
 - Credit scores from different scoring models
- **Trade Lines (Accounts):**
 - Account details: Creditor, type, opened/closed dates
 - Credit terms: Limit, payment amount, loan amount
 - Current status: Balance, utilization, payment status
 - Payment history: 24-84 months of payment records

Credit Report Inquiries

- **Hard Inquiries:**
 - Result from credit applications
 - Visible to all lenders
 - Impact credit scores
- **Soft Inquiries:**
 - Account reviews, promotional screening
 - Only visible to consumer
 - No impact on credit scores

Credit Scores

- Two major scoring models:
 - FICO
 - VantageScore
- Credit Score Ranges:
 - Excellent: 800+
 - Very Good: 740-799
 - Good: 670-739
 - Fair: 580-669
 - Poor: Below 580

Prescreen Campaigns to Acquire New Clients



Business Problem

- Bank wants to acquire new credit card customers through targeted marketing
- Marketing budget is limited, so need to focus on high-response rate prospects
- Must comply with regulatory requirements (FCRA, ECOA)
- Need to balance response rates against credit risk
- Campaign costs include: mailing, credit bureau fees, and expected credit losses

Business Solution

- Develop data-driven selection criteria for credit bureau prescreening
- Selection criteria typically includes:
 - Credit scores from major bureaus (FICO, VantageScore)
 - Geographic targeting based on market strategy
 - Key bureau attributes (payment history, utilization, delinquencies)
 - Existing banking relationship data if available
 - Historical campaign response patterns
 - Age of credit history and recent inquiries
- Establish control groups to measure campaign effectiveness

Machine Learning Problem

- Requires historical campaign data, response flags, and relevant bureau attributes
- Predict which prospects are most likely to accept the offer
- Develop a model that maximizes AUC and estimate a probability threshold that maximizes the F1 score, precision, or recall

Machine Learning Solution

- Train a classification model on labeled data to identify strong responders
- Incorporate performance metrics such as AUC, F1 score, lift chart, and gains chart to gauge effectiveness
- Score prospects, select high-scoring individuals for the campaign, and measure the lift and gains chart over a control group

Prescreening Compliance Framework

- Fair Credit Reporting Act (FCRA)
 - Verify permissible purpose and consumer opt-out status prior to prescreening
 - Ensure firm offers match the criteria used during prescreening
 - Include clear opt-out notices and instructions in every prescreened offer
 - If declined after prescreening, provide an Adverse Action Notice
- Equal Credit Opportunity Act (ECOA)
 - Prohibits discrimination based on race, color, religion, national origin, sex, marital status, or age
 - Apply consistent underwriting standards to all applicants
 - Maintain documentation to demonstrate non-discriminatory practices and facilitate regulatory reviews

Citations

- Azad, Rohan. 2024. “Leading Indian Bank & CIBIL Real-World Dataset.”
<https://www.kaggle.com/datasets/saurabhbadole/leading-indian-bank-and-cibil-real-world-dataset/data>.
- Fannie Mae. 2024. “Single-Family Loan Performance Data.”
<https://capitalmarkets.fanniemae.com/credit-risk-transfer/single-family-credit-risk-transfer/fannie-mae-single-family-loan-performance-data>.
- George, Nate. 2019. “All Lending Club Loan Data.”
<https://www.kaggle.com/datasets/wordsforthewise/lending-club>.
- Le Borgne, Yann-Aël, Wissam Siblini, Bertrand Lebichot, and Gianluca Bontempi. 2022. *Reproducible Machine Learning for Credit Card Fraud Detection - Practical Handbook*. Université Libre de Bruxelles. <https://github.com/Fraud-Detection-Handbook/fraud-detection-handbook>.